

Eegapuri Lokeshwar Reddy

Data Engineer

+91 9959242262 ◇ elokesh4292@gmail.com ◇ Telangana, India ◇ [LinkedIn](#) ◇ [GitHub](#) ◇ [Website](#)

PROFESSIONAL SUMMARY

Data Engineer with 3 years of experience designing and supporting batch and streaming data solutions, ETL/ELT workflows, and enterprise data warehousing across production and disaster-recovery environments. Hands-on expertise in dimensional modeling, OLAP systems, and data marts, gained through extensive EDW experience. Proficient in relational database architectures, big-data frameworks, and cloud-based analytics, with a strong interest in data lake and lakehouse technologies for unified platform solutions. Combines deep operational discipline with a problem-solving mindset to build robust, real-time and batch data solutions.

SKILLS

Databases & Warehousing: Teradata, PostgreSQL, Data Warehousing, Data Modeling, Dimensional Modeling, ERD, Data Lakes, Data Lakehouse, Delta Lake

ETL / Workflow Orchestration: Ab Initio, Apache Airflow, Shell Scripting

Big Data & Batch Processing: Apache Spark (SparkSQL, PySpark), Hadoop (HDFS, YARN)

Streaming & Messaging: Apache Kafka, Confluent Kafka, ksqlDB

Cloud & DevOps Tools: AWS Glue, AWS S3, Databricks, Red Hat Linux, Git, Looker Studio

Programming Languages & Libraries: Python, SQL, Pandas, NumPy, SQLAlchemy, DB API, BeautifulSoup

Collaboration & Issue Tracking: JIRA, Remedy

Soft Skills: Problem Solving, Communication, Analytical Thinking, Adaptability, Attention to Detail, Risk Management, Decision-Making, Emotional discipline, Patience, Reliability, Knowledge Sharing

CERTIFICATIONS

Confluent Fundamentals Accreditation, Confluent

Academy Accreditation - Databricks Fundamentals, Databricks

IBM Data Engineering Professional Certificate, IBM (In Progress)

PERSONAL PROJECTS

RentLok: Real-Time Rental Management System | [Link](#)

- **Tech Stack:** Android, FastAPI, PostgreSQL, Apache Kafka, Kafka Connect, ksqlDB, SQLAlchemy, Pydantic, JSON
 - Built an end-to-end **OLTP** data pipeline for a rental management system with a native **Android app** and backend powered by **FastAPI**, **PostgreSQL**, and **Apache Kafka**.
 - Developed secure RESTful **APIs** over **HTTPS** to process **JSON** data, enforced schema validation with **Pydantic**, and used Kafka Connect to stream PostgreSQL data into Kafka topics.
 - Implemented **ksqlDB** for real-time transformations and aggregations, enabling live in-app metrics; optimized the **PostgreSQL** schema for low-latency access and modular microservice deployment.
- **Key Learning:**
Successfully transitioned from OLAP to OLTP system design, demonstrating the ability to build reliable, API-driven, real-time data pipelines and translate business requirements into scalable, end-to-end data solutions.

Telecom Data Pipeline | [Link](#)

- **Tech Stack:** Apache Airflow, Apache Spark (PySpark), PostgreSQL, Shell Scripting, Python
 - Designed a batch processing **ETL pipeline** to process millions of CDR records using **Airflow** for orchestration and **PySpark** for scalable, distributed transformations.
 - Developed fraud detection logic to identify suspicious patterns like high-frequency calls and short-duration international interactions.
 - Implemented audit logging and error tracking, improving pipeline observability and operational reliability.
- **Key Learning:**
Demonstrated the ability to design scalable, distributed data transformations using PySpark, while enhancing reliability through fraud detection logic and robust logging.

E-Commerce Data Pipeline | [Link](#)

- **Tech Stack:** Apache Airflow, MySQL, PostgreSQL, Foreign Data Wrapper (FDW), Python, Pandas, SQL
 - Built a modular batch ETL pipeline using Airflow to automate CSV ingestion and transformations across MySQL and PostgreSQL.
 - Enabled cross-database querying using PostgreSQL FDW, secured credentials with `.my.cnf` and `.pgpass`, and optimized performance through indexing.
 - Exported cleaned data for downstream analysis, showcasing a reusable and scalable data integration workflow.

- **Key Learning:**
Strengthened skills in Airflow orchestration, relational database data ingestion, and secure, cross-database batch processing aligned with real-world data engineering practices.

Game Development & 3D Modeling | [Link](#)

- **Tech Stack:** Unity, C#, Blender
 - Developed a PC game (“**Cubon**”) using **Unity** and **C#**, implementing player controls, interactive mechanics, and UI features.
 - Created and animated 3D models in **Blender**, applied texturing and rendering, and integrated assets into the Unity engine.
- **Key Learning:**
Gained hands-on experience in software development, event-driven programming, and 3D rendering pipelines — combining creative and technical skills foundational to full-cycle product development.

WORK EXPERIENCE

System Engineer
Tata Consultancy Service (TCS)

Apr '19 — Feb '22
Hyderabad

Project 1: Saudi Telecom Company(STC) – Enterprise Data Warehouse (EDW) Operations

Environment: Production & Disaster Recovery (DR)

Tools: Ab Initio, Teradata, Red Hat Linux, SQL, Shell Scripting, Jira, Remedy

- Monitored **2000+** Enterprise-Scale **ETL Jobs** across **OLAP** and data mart environments, responding to daily requests and escalations to ensure uninterrupted data flow.
- Oversaw ingestion from **CRM, CDR, and IRB** systems via **SFTP** and **SQL** Extraction, automating workflows to reduce manual interventions.
- Handled **200+ Change Requests**, including deployment of **45+ New ETL Jobs**, optimizations, data reloads, and decommissions—preventing multiple potential **SLA** breaches.
- Developed and optimized **SQL Queries** to track business-critical **KPIs** (postpaid, prepaid, landline, sales) and automated discrepancy reporting for proactive resolution.
- Created **Shell Scripts** to monitor system resources, detect transformation failures, and clean unwanted files—recovering **30% Storage space**, improving operational efficiency by **35%**, and ensuring **100% Data Completeness**.
- Collaborated in **4–5 DR Drills** to validate failover processes: tested job workflows, source connections, and scheduling configurations, and maintained documentation for seamless DR execution.

Project 2: Saudi Telecom Company(STC) – Kafka-Based Event Streaming Operations. (May '21 - Feb '22)

Environment: Production

Tools: Confluent Kafka, Python, Kafka Client Libraries, ksqlDB, Avro, JSON

- Managed operations and maintenance of **250+ Kafka Topics**, supporting real-time event streaming and **ksqlDB** stream processing alongside EDW batch workflows.
- Ensured schema compatibility by handling **Avro** and **JSON** message formats and resolving connector issues (JDBC/ODBC, HTTP integrations, SSL errors, schema mismatches).
- Developed **Python** Scripts using Kafka client libraries to automate health checks, monitor consumer lag, and send real-time alerts—significantly reducing manual monitoring effort.
- Addressed core **Kafka** issues, including **RBAC** Permission errors, driver mismatches, performance tuning, and schema evolution conflicts to maintain reliable data delivery.

EDUCATION

B-Tech in Mechanical Engineering, CVR College of Engineering (GPA: 9.56)

2015 — 2019
Hyderabad

Intermediate in MPC, Narayana Junior College (GPA: 9.7)

2013 — 2015
Hyderabad

ACHIEVEMENTS & AWARDS

- On the **Spot Award** for outstanding performance and exceptional contributions to **ETL operations** and **system optimization**.
- Received **client appreciation** for ensuring seamless system operations and maintaining **100% KPI availability** with **zero SLA breaches**.
- Cleared the first round of **TCS CodeVita Season 7 (2018)**, a global coding competition, demonstrating strong programming and problem-solving skills.